

```

### Importing the dataset - in the form of batches - using dataloader package - mini batch descent - size =64
import torch
from torchvision import datasets, transforms
from torch.utils.data import DataLoader

import matplotlib.pyplot as plt
import numpy as np
transform = transforms.Compose([transforms.ToTensor(),
transforms.Normalize((0.5,), (0.5,))])
train_dataset = datasets.MNIST(root='./data', train=True,
transform=transform, download=True)
test_dataset = datasets.MNIST(root='./data', train=False,
transform=transform, download=True)
train_loader = DataLoader(dataset=train_dataset, batch_size=64,
shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=64,
shuffle=False)

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

class MLP_torch(nn.Module):
    def __init__(self, input_layer_size=784,
first_hidden_layer_size=500,
second_hidden_layer_size=250,
third_hidden_layer_size=100,
output_layer_size=10, regulatisation_constant =
0.0001):
        super(MLP_torch, self).__init__()
        self.regularisation_constant = regulatisation_constant
        self.first_layer = nn.Linear(input_layer_size,
first_hidden_layer_size)
        self.second_layer = nn.Linear(first_hidden_layer_size,
second_hidden_layer_size)
        self.third_layer = nn.Linear(second_hidden_layer_size,
third_hidden_layer_size)
        self.output_layer = nn.Linear(third_hidden_layer_size,
output_layer_size)

    def forward(self, x):
        x = torch.sigmoid(self.first_layer(x))
        x = torch.sigmoid(self.second_layer(x))
        x = torch.sigmoid(self.third_layer(x))
        x = self.output_layer(x)
        x = torch.softmax(x, dim=1)
        return x

```

```

def entropy_class_loss(self, Y, Y_hat):
    epsilon = 1e-10
    Y_hat = torch.clamp(Y_hat, epsilon, 1 - epsilon)

    loss = -torch.sum(Y * torch.log(Y_hat)) / Y.size(0)

    l2_norm_sum = 0
    for param in self.parameters():
        param_norm = torch.norm(param, p=2)
        l2_norm_sum += param_norm

    return loss + self.regularisation_constant*l2_norm_sum

import torch
import torch.nn as nn
import matplotlib.pyplot as plt
from torch.optim import Adam
import torch.nn.functional as F
import torch.nn.functional as F

def train(model, optimizer, train_loader, test_loader, epochs,
plot_interval=200, device='cpu'):
    training_loss_list = []
    test_loss_list = []
    iteration_counts = []
    current_iteration = 0

    for epoch in range(epochs):
        model.train()
        running_loss = 0
        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.view(images.size(0), -
1).to(device), labels.to(device)
            labels_one_hot = F.one_hot(labels,
num_classes=10).float().to(device)
            optimizer.zero_grad()
            outputs = model(images)
            loss = model.entropy_class_loss(outputs, labels_one_hot)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            current_iteration += 1

            if current_iteration % plot_interval == 0:
                iteration_counts.append(current_iteration)
                avg_train_loss = running_loss / (batch_idx + 1)
                training_loss_list.append(avg_train_loss)

```

```

        model.eval()
        total_test_loss = 0
        total_test_samples = 0
        with torch.no_grad():
            for test_images, test_labels in test_loader:
                test_images, test_labels =
test_images.view(test_images.size(0), -1).to(device),
test_labels.to(device)

                test_outputs = model(test_images)
                test_labels_one_hot = F.one_hot(test_labels,
num_classes=10).float().to(device)
                test_loss =
model.entropy_class_loss(test_outputs, test_labels_one_hot)
                total_test_loss += test_loss.item() *
test_images.size(0)
                total_test_samples += test_images.size(0)

        avg_test_loss = total_test_loss / total_test_samples
        test_loss_list.append(avg_test_loss)
        print(f'Current Iteration {current_iteration},
Training Loss: {avg_train_loss:.4f}, Test Loss: {avg_test_loss:.4f}')

        print(f'Epoch {epoch + 1}/{epochs} ')

    plt.figure(figsize=(20, 17))
    plt.plot(iteration_counts, training_loss_list, label='Training
Losses')
    plt.plot(iteration_counts, test_loss_list, label='Test Losses')
    plt.xlabel('Iterations')
    plt.ylabel('Loss')
    plt.legend()
    plt.grid(True)
    plt.show()


model = MLP_torch()
optimizer = Adam(model.parameters())
train(model, optimizer, train_loader, test_loader, epochs=15)

Current Iteration 200, Training Loss: 14.8370, Test Loss: 10.3845
Current Iteration 400, Training Loss: 11.9816, Test Loss: 8.0058
Current Iteration 600, Training Loss: 10.6772, Test Loss: 7.9521
Current Iteration 800, Training Loss: 9.9824, Test Loss: 7.6136
Epoch 1/15
Current Iteration 1000, Training Loss: 7.7351, Test Loss: 7.5165
Current Iteration 1200, Training Loss: 7.5839, Test Loss: 7.5872
Current Iteration 1400, Training Loss: 7.6201, Test Loss: 7.6701
Current Iteration 1600, Training Loss: 7.6275, Test Loss: 7.5003
Current Iteration 1800, Training Loss: 7.6283, Test Loss: 7.4399
Epoch 2/15

```

Current Iteration	2000,	Training Loss:	5.2271,	Test Loss:	4.3022
Current Iteration	2200,	Training Loss:	4.5080,	Test Loss:	4.2277
Current Iteration	2400,	Training Loss:	4.3342,	Test Loss:	4.1028
Current Iteration	2600,	Training Loss:	4.0838,	Test Loss:	2.5926
Current Iteration	2800,	Training Loss:	3.6474,	Test Loss:	2.0018
Epoch 3/15					
Current Iteration	3000,	Training Loss:	1.6120,	Test Loss:	1.6286
Current Iteration	3200,	Training Loss:	1.5974,	Test Loss:	1.7613
Current Iteration	3400,	Training Loss:	1.5930,	Test Loss:	1.3964
Current Iteration	3600,	Training Loss:	1.5685,	Test Loss:	1.5625
Epoch 4/15					
Current Iteration	3800,	Training Loss:	1.4752,	Test Loss:	1.4325
Current Iteration	4000,	Training Loss:	1.4293,	Test Loss:	1.6935
Current Iteration	4200,	Training Loss:	1.3398,	Test Loss:	1.6296
Current Iteration	4400,	Training Loss:	1.3324,	Test Loss:	1.2966
Current Iteration	4600,	Training Loss:	1.3060,	Test Loss:	1.4562
Epoch 5/15					
Current Iteration	4800,	Training Loss:	1.2556,	Test Loss:	1.1559
Current Iteration	5000,	Training Loss:	1.1453,	Test Loss:	1.1667
Current Iteration	5200,	Training Loss:	1.1771,	Test Loss:	1.0848
Current Iteration	5400,	Training Loss:	1.1882,	Test Loss:	1.3754
Current Iteration	5600,	Training Loss:	1.1598,	Test Loss:	1.2138
Epoch 6/15					
Current Iteration	5800,	Training Loss:	1.0559,	Test Loss:	1.1979
Current Iteration	6000,	Training Loss:	1.0320,	Test Loss:	1.0801
Current Iteration	6200,	Training Loss:	1.0326,	Test Loss:	1.0596
Current Iteration	6400,	Training Loss:	1.0306,	Test Loss:	1.1444
Epoch 7/15					
Current Iteration	6600,	Training Loss:	1.0287,	Test Loss:	1.0559
Current Iteration	6800,	Training Loss:	0.9960,	Test Loss:	1.1091
Current Iteration	7000,	Training Loss:	1.0020,	Test Loss:	1.2430
Current Iteration	7200,	Training Loss:	1.0069,	Test Loss:	0.9973
Current Iteration	7400,	Training Loss:	1.0136,	Test Loss:	1.1068
Epoch 8/15					
Current Iteration	7600,	Training Loss:	1.0113,	Test Loss:	1.1766
Current Iteration	7800,	Training Loss:	0.9947,	Test Loss:	1.0521
Current Iteration	8000,	Training Loss:	0.9866,	Test Loss:	1.0694
Current Iteration	8200,	Training Loss:	0.9803,	Test Loss:	1.1300
Current Iteration	8400,	Training Loss:	0.9800,	Test Loss:	1.0557
Epoch 9/15					
Current Iteration	8600,	Training Loss:	0.8429,	Test Loss:	0.9627
Current Iteration	8800,	Training Loss:	0.9082,	Test Loss:	1.3135
Current Iteration	9000,	Training Loss:	0.9263,	Test Loss:	1.2462
Current Iteration	9200,	Training Loss:	0.9679,	Test Loss:	0.9925
Epoch 10/15					
Current Iteration	9400,	Training Loss:	0.8419,	Test Loss:	0.8898
Current Iteration	9600,	Training Loss:	0.8386,	Test Loss:	0.9911
Current Iteration	9800,	Training Loss:	0.8898,	Test Loss:	1.0584
Current Iteration	10000,	Training Loss:	0.8912,	Test Loss:	1.1633

Current Iteration 10200, Training Loss: 0.9041, Test Loss: 0.9412
Epoch 11/15
Current Iteration 10400, Training Loss: 0.7138, Test Loss: 1.0924
Current Iteration 10600, Training Loss: 0.8758, Test Loss: 0.8545
Current Iteration 10800, Training Loss: 0.8573, Test Loss: 1.7867
Current Iteration 11000, Training Loss: 0.8682, Test Loss: 0.9807
Current Iteration 11200, Training Loss: 0.8700, Test Loss: 0.9192
Epoch 12/15
Current Iteration 11400, Training Loss: 0.7863, Test Loss: 0.9207
Current Iteration 11600, Training Loss: 0.8303, Test Loss: 1.0136
Current Iteration 11800, Training Loss: 0.8459, Test Loss: 1.0285
Current Iteration 12000, Training Loss: 0.8274, Test Loss: 0.8618
Epoch 13/15
Current Iteration 12200, Training Loss: 0.7547, Test Loss: 1.1169
Current Iteration 12400, Training Loss: 0.8412, Test Loss: 0.9867
Current Iteration 12600, Training Loss: 0.8395, Test Loss: 0.7782
Current Iteration 12800, Training Loss: 0.8365, Test Loss: 1.0042
Current Iteration 13000, Training Loss: 0.8315, Test Loss: 0.8685
Epoch 14/15
Current Iteration 13200, Training Loss: 0.7687, Test Loss: 0.8922
Current Iteration 13400, Training Loss: 0.7642, Test Loss: 0.9481
Current Iteration 13600, Training Loss: 0.7562, Test Loss: 0.8029
Current Iteration 13800, Training Loss: 0.7733, Test Loss: 0.8487
Current Iteration 14000, Training Loss: 0.7834, Test Loss: 0.8111
Epoch 15/15



```
import torch
import numpy as np
from sklearn.metrics import confusion_matrix, accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt

def final_test(model, test_loader):
    model.eval()
    all_preds = []
    all_labels = []
    with torch.no_grad():
        for images, labels in test_loader:
            images = images.view(images.size(0), -1)
            outputs = model(images)
            _, preds = torch.max(outputs, 1)
            all_preds.append(preds.cpu().numpy())
            all_labels.append(labels.cpu().numpy())
    all_preds = np.concatenate(all_preds)
```


```

all_labels = np.concatenate(all_labels)
accuracy = accuracy_score(all_labels, all_preds)
print(f'Final Test Accuracy: {accuracy * 100:.2f}%')
conf_matrix = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(20, 17))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

```
final_test(model, test_loader)
```

Final Test Accuracy: 96.27%



```
final_test(model, train_loader)
```

Final Test Accuracy: 97.17%

